

CITS1402 Week 6 Lab

Gordon Royle

2022 Semester Two

This is an *individual* assessment, and so the query that you submit should be devised by you alone, typed into the submission file by you alone, tested by you alone, and then submitted entirely by yourself. You may discuss the assignment with the lab facilitators and your fellow students in *general terms only*. If you are unsure of the line between legitimate discussion and forbidden collaboration, then you can apply the following simple test: If the discussion is purely verbal and mentions no SQLite code, then it is probably legitimate. On the other hand, if you describe actual SQLite code in your discussion, or you show someone your screen, or email someone your files, then it is not legitimate. So you can say “I think you need to figure out which tables have the information, and then join them”, but nothing more detailed. If someone asks you to *look* at their code, or to *show* them your code, then you should politely refuse.

You should submit your solutions to the selected lab questions at `secure.csse.uwa.edu.au/run/cssubmit`. The selected questions are highlighted by a **bright yellow box** containing the word `cssubmit` and then a filename.

For each of the highlighted questions, you should save your SQL query into the specified filename and then upload the resulting file (saved as plain text) to `cssubmit`. The validation system will prevent you from submitting files with incorrect names.

Make sure that you *test* your queries by

- Typing `cat A1.sql` at the Powershell/Terminal prompt to view the contents of the file, making sure that there are no formatting commands or strange characters.
- Using `sqlite3` to open the sample database, and then `.read A1.sql` at the `sqlite3` prompt, to see the output that I (my marking scripts) will see.
- Double-check the specifications to ensure that the output has the right columns, in the correct order (column order matters), and that the output rows are correct and complete.
- Double check that you have actually *submitted* the files to `cssubmit` and not just *validated* them.

Of course, replace `A1.sql` by whichever file you are testing.

Use the filename that is *in the bright yellow box*, no matter where the question occurs in the list of questions.

Learning Aims

This lab is primarily concerned with aggregation and summary functions, therefore using `GROUP BY` and `HAVING`.

We will use the databases from previous labs.

For this lab you may find it easier to write (and hence read) your queries if you give the tables that you use short aliases. For example, to count the number of employees located in each office of Classic Models, we could write:

```
SELECT O.city, COUNT(*) as numEmployees
FROM offices O JOIN employees USING (officeCode)
GROUP BY O.officeCode, O.city;
```

The FROM statement assigns the table alias **O** to the table **offices** and we can then use this shortcut anywhere else, including the

Questions

1. Using the IMDB Top-250, write a single SQL query that lists, for each *year* represented in the database (this is the column **premiered**) the *number of movies* from that year in the IMDB Top-250.

SELF-CHECK HINT There are six movies from 1957 that are still in the Top-250.

2. Using Classic Models, write a single SQL query that lists, for each city, the *name* of the city, and the *number of employees* based in that city.

SELF-CHECK HINT There are four employees in Sydney.

3. Using **AustOpen.db**, write a single SQL query that for each year in the database lists the *year* and the *number of matches* won by Australian players in the Women's Singles for that year. (Remember that the data for Women's Singles is in the table **WTAResult**.)

SELF-CHECK HINT 2022 was a bumper year, because Australian women won a total of 10 matches.

4. **Submission: A8.sql** Using **AustOpen.db**, write a single SQL query lists the *year*, *player name* and *number of matches won* for each of the Australian players in the Women's Singles. (You need not consider players who did not win any matches — this substantially simplifies the query because you only need to consider names in the **winnerName** column.)

SELF-CHECK HINT In 2022, Ashleigh Barty was the Women's Champion, which she achieved by winning the maximum of 7 matches, so make sure that one of your rows is (2022, Ashleigh Barty, 7).

5. Using Classic Models, write a single SQL query that lists, for each product line, the *product line* and the *price paid* by Classic Models of the most expensive product in that product line. Recall that the price paid by Classic Models is (Note that you do not have the list the actual product, just its price.)

SELF-CHECK HINT The most expensive product in the **Ships** line costs \$82.34, so one of the rows of the output table should be:

Ships 82.34

6. **Submission: A9.sql** Using Classic Models, write a single SQL query that lists, for each order, its *order number*, the *number of line-items* in the order, and the *total cost* of the order, rounded to two decimal places (look up and use the **ROUND** function for this). The output table should have three columns, and should list more expensive orders before cheaper ones.

Avoid using more tables than you need.

SELF-CHECK HINT Order 10165 contains 18 line-items and has a total cost of 67392.85.

10100	Vintage Cars	4	10223.83
10101	Vintage Cars	4	10549.01
10102	Vintage Cars	2	5494.78
10103	Classic Cars	3	14548.88
10103	Trucks and B	7	20987.03
10103	Vintage Cars	6	14683.04
10104	Classic Cars	7	22003.45
10104	Trains	2	4476.87
10104	Trucks and B	4	13725.88

Figure 1: Orders broken down by product lines

7. We want to find out whether *younger players* have an advantage in the Australian Open Men's Singles. Write a single SQL query that calculates the percentage of matches won by the younger player. For each year, the output table should list the *year* and the *percentage of matches* won by the younger player, with the percentage rounded to 2 decimal places.

Every tournament has a total of 127 matches and this is a value that is very unlikely to change, so you may directly hard-code the value 127 into your query.

SELF-CHECK HINT In 2001, 46.46 % of the matches were won by the younger of the two players.

8. **Submission: A10.sql** Do some male Australian Open players take significantly longer than others to win the matches that they do win? Write a single SQL query that lists, for each winner, the *name* of the player, the *number of matches* they have won, and the *average length* of those matches, rounded to 2 decimal places. Restrict the output to players who have won at least 10 matches, and sort the rows so that the fastest players are at the top. Note that we are only considering the matches that a player wins and remember that "at least 10" includes 10.

SELF CHECK HINT Andre Agassi is the fastest in this list, having won 30 matches with an average duration of just 108.13 minutes per match.

9. Using Classic Models, write a single SQL query that lists, for each combination of order and product line, the *order number*, the *product line*, the *number* of line-items of that product line in the order, and the *total cost* of those line items, rounded to 2 decimal places. (You only need to list the product lines that actually occur in the order.)

SELF-CHECK HINT: Figure 1 shows that Order 10103 contains 3 line-items of Classic Cars, 7 of Trucks and Buses and 6 of Vintage Cars. The three line-items of Classic Cars cost a total of 14548.88.

10. **Submission: A11.sql** Using Classic Models, write a single SQL query that lists each *order number* along with the *net profit*, rounded to 2 decimal places, that Classic Models made from that order. The table should have two columns named *orderNumber* and *netProfit*.

SELF-CHECK HINT: Order 10422 has two line items only, so is a good choice to check your code. This order consists of 51 units of S18_1342, which were sold for \$91.44 each, and 25 units of S18_1367 which were sold for \$47.44 each, for a total cost of \$5849.44. As Classic Models paid \$60.62 for each S18_1342 and \$24.26 for each S18_1367, the items in this order cost Classic Models a total of \$3698.12. Therefore the profit margin on this particular order is \$5849.44-\$3698.12 = \$2151.32.

11. Classic Models has a "Important Customer" list, which is a list of all the customers who have made orders that (added together) have a total cost of at least \$500000 (half a million dollars). Write a single SQL query that lists the *customer name* and *total order value* of these Important Customers.

LEARNING BREAK

The tables `WTAResult` and `ATPResult` contain data about individual matches in a different way to `AFLResult`,

In `WTAResult` and `ATPResult` it is easy to identify the winner, because the winner's name is always in the same column, namely `winnerName`. But in `AFLResult` the name of the winning team is sometimes in the column `homeTeam` and sometimes in the column `awayTeam`, and of course sometimes there is no winner at all, because the game can end in a draw.

So for `AFLResult` it is not clear how to write a `SELECT` statement that chooses the winning team — it requires some sort of *conditional execution* that changes the column being selected according to the values in other columns. In a procedural programming language, this would be easy to write, with something like the following:

```
if homeScore > awayScore:
    return homeTeam
elif homeScore < awayScore:
    return awayTeam
else:
    return 'DRAW'
```

But SQL is a declarative language, and hence has no direct “control flow” statements. Instead there are a couple of special commands that permit the user some limited control to alter the value returned based on a boolean condition.

The first of these is the `CASE` expression which has the following syntax:

```
CASE WHEN <boolean_1> THEN <expr_1>
      WHEN <boolean_2> THEN <expr_2>
      WHEN ...
      ELSE <expr_3>
END
```

This entire expression is used as part of the `SELECT` statement to specify *one of the columns* of the output table. When a row is processed, the value of this column is the *first expression* for which the associated boolean condition is true. The ellipsis (dots) just indicate that you can use as many `WHEN` clauses as necessary.

An example makes this clearer:

```
SELECT year,
       round,
       CASE WHEN (homeScore > awayScore) THEN homeTeam
            WHEN (homeScore < awayScore) THEN awayTeam
            ELSE NULL
       END
FROM AFLResult;
```

This code defines a table with *three columns* — the first column is called `year` and is just copied from the row being processed, the second column is called `round` and is just copied from the row being processed, while the third column is *calculated* according to the specifications of the `CASE` statement. This compares `homeScore` to `awayScore` and chooses either `homeTeam`, `awayTeam` or `NULL` according to whether `homeScore` is greater than, less than, or equal to `awayScore`.

You can rename the column defined by the `CASE` statement in the usual way.

```

SELECT year,
       round,
       CASE WHEN (homeScore > awayScore) THEN homeTeam
            WHEN (homeScore < awayScore) THEN awayTeam
            ELSE NULL
       END AS winner
FROM AFLResult;

```

There is a second way to do this that has only recently been added to SQLite (version 3.32 or higher), namely the `IIF` statement, which is perhaps easier to read when there are only two choices.

`IIF` was introduced in SQLite 3.32, so will not work if you are using an earlier version of SQLite. You can check which version you have if you use the dot command `.version` in the `sqlite3` terminal program, and you can install a more recent version if you wish.

This code lists the *older player* from each match, by simply testing whether `winnerAge` is greater than `loserAge` or not. (This ignores the possibility that both players have exactly the same age.)

```

SELECT yr,
       IIF(winnerAge > loserAge, winnerName, loserName) as olderPlayer
FROM ATPResult;

```

12. Using `AFLResult.db`, use a single SQL query to produce a list of the winning scores for each game. Each row of the 3-column table should list the *year*, the *round* and the *winning score*. (If the match was a draw, then we view either team's score as the winning score).

SELF-CHECK HINT: One of the games in Round 1, 2012 was won by a team scoring exactly 100 points. Make sure this information appears in your output.

13. Using `AFLResult.db`, use a single SQL query to find the *average winning score* for each of the years in the database. Your output table should have two columns, one for the *year* and another for the *average winning score* for all the matches in that year.

SELF-CHECK HINT: The average winning score in 2021 was 94.55050505050505.

14. (Challenge Question) Using `AFLResult.db`, use a single SQL query to find the *average winning score* for *each team* in *each year* in the database. Your output table should have three columns, the first for the *year*, the second for the *team*, and the third for the *average winning score* made by that team in all the games that they won in that particular year. In this question, *do not count draws* as a win for either team.

For example in 2018, Carlton only won two matches. They beat Essendon 91-78 in Round 8, and beat Gold Coast Suns 79-44 in Round 18. Therefore they have two wins, and an average winning score of $(79 + 91)/2 = 85$.